

# Model Training and Evaluating it

Thursday 11<sup>th</sup> December, 2025



We look to implement and compare five different forecasting models.

```
#Setting everything up
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import os
import sys

from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.preprocessing import StandardScaler
from statsmodels.tsa.arima.model import ARIMA
import tensorflow as tf
from tensorflow.keras import layers, models

sys.path.append(os.path.abspath(os.path.join('..'))))
from scripts.data_process import make_supervised_frame, time_series_split
from scripts.data_modeling import compute_regression_metrics, directional_accuracy, make_lst

np.random.seed(42)
tf.random.set_seed(42)
plt.rcParams["figure.figsize"] = (10, 5)
plt.rcParams["axes.grid"] = True

2025 -12 -11 01:53:06.154772: I external/local_xla/xla/tsl/cuda/cudart_stub.cc:31] Could not
2025 -12 -11 01:53:06.253614: I tensorflow/core/util/port.cc:153] oneDNN custom operations a
2025 -12 -11 01:53:09.995778: I tensorflow/core/platform/cpu_feature_guard.cc:210] This Tens
To enable the following instructions: AVX2 AVX512F AVX512_VNNI FMA, in other operations, reb
2025 -12 -11 01:53:20.679920: I tensorflow/core/util/port.cc:153] oneDNN custom operations a
2025 -12 -11 01:53:20.692153: I external/local_xla/xla/tsl/cuda/cudart_stub.cc:31] Could not

#Loading data
df = pd.read_csv("../data/sp500.csv", index_col = "Date", parse_dates = True)
```

```

df_sup = make_supervised_frame(df, target_col="LogReturn", horizon=1, lags=3)

df_sup.head()

```

Date	Clos	Hig	Low	Open	Volum	Ma1	Ma5	Ma10	Ma50	logRet	Min	Max	MA10D	LogReg	Reg	Reglag	lag3
1990-03-13	36.889	37.035	35.985	35.985	0.000000	0.000000	0.000000	0.000000	0.000000	-0.005620	0.002582	0.002187	0.009805	0.007915	0.212164	0.006901	
1990-03-14	36.887	37.035	35.985	35.985	0.000000	0.000000	0.000000	0.000000	0.000000	-0.007897	0.002582	0.002187	0.009805	0.007915	0.006901		
1990-03-15	38.070	37.035	35.985	35.985	0.000000	0.000000	0.000000	0.000000	0.000000	-0.007897	0.002582	0.002187	0.009805	0.007915	0.006901		
1990-03-16	38.070	37.035	35.985	35.985	0.000000	0.000000	0.000000	0.000000	0.000000	-0.007897	0.002582	0.002187	0.009805	0.007915	0.006901		
1990-03-19	41.940	39.035	39.035	39.035	0.000000	0.000000	0.000000	0.000000	0.000000	-0.007897	0.002582	0.002187	0.009805	0.007915	0.006901		

```

df_train, df_val, df_test = time_series_split(df_sup, train_frac=0.6, val_frac=0.2)

feature_cols = [c for c in df_sup.columns if c != "y"]

X_train, y_train = df_train[feature_cols], df_train["y"]
X_val, y_val = df_val[feature_cols], df_val["y"]
X_test, y_test = df_test[feature_cols], df_test["y"]

print("Train size:", len(X_train))
print("Val size:", len(X_val))
print("Test size:", len(X_test))

Train size: 5401
Val size: 1800
Test size: 1801

#baseline
y_test_naive = df_test["LogReturn_lag1"]

baseline_metrics = compute_regression_metrics(y_test, y_test_naive)
baseline_dir_acc = directional_accuracy(y_test, y_test_naive)

```

```

print("Naive baseline metrics:", baseline_metrics)
print("Naive baseline directional accuracy:", baseline_dir_acc)

Naive baseline metrics: {'RMSE': np.float64(0.016940516181126818), 'MAE': 0.0116095086299034}
Naive baseline directional accuracy: 0.5091615769017213

```

## 1 Linear Baseline

After the naive baseline, we use a Linear Regression model. According to the Efficient Market Hypothesis (EMH), daily returns should be a random walk, meaning our linear model should have a  $R^2$  close to zero. This will be a good indicator of whether complex models can extract any signal.

```

# Linear Regression
linreg = LinearRegression()
linreg.fit(X_train, y_train)

y_val_lr = linreg.predict(X_val)
y_test_lr = linreg.predict(X_test)


lr_val_metrics = compute_regression_metrics(y_val, y_val_lr)
lr_test_metrics = compute_regression_metrics(y_test, y_test_lr)
lr_dir_acc = directional_accuracy(y_test, y_test_lr)

print("Linear Regression - Validation Metrics:", lr_val_metrics)
print("Linear Regression - Test Metrics:", lr_test_metrics)
print(f"Linear Regression - Directional Accuracy: {lr_dir_acc:.2%}")

Linear Regression - Validation Metrics: {'RMSE': np.float64(0.008868328927820853), 'MAE': 0.008868328927820853}
Linear Regression - Test Metrics: {'RMSE': np.float64(0.013779692199568672), 'MAE': 0.010000000000000002}
Linear Regression - Directional Accuracy: 46.42%

plt.scatter(y_test, y_test_lr, alpha=0.5)
plt.xlabel("Actual LogReturn")
plt.ylabel("Predicted LogReturn (LR)")
plt.title("Linear Regression - Test Predictions vs Actual")
plt.savefig("../images/LR.png")
plt.show()

```



The Linear Regression model had an  $R^2$  score of -0.17. So, this implies the model performs worse than just guessing the average return.

```

rf = RandomForestRegressor(n_estimators=200, max_depth=5, random_state=42)
rf.fit(X_train, y_train)

y_val_rf = rf.predict(X_val)
y_test_rf = rf.predict(X_test)

rf_val_metrics = compute_regression_metrics(y_val, y_val_rf)
rf_test_metrics = compute_regression_metrics(y_test, y_test_rf)
rf_dir_acc = directional_accuracy(y_test, y_test_rf)

print("Random Forest - validation metrics:", rf_val_metrics)
print("Random Forest - test metrics:", rf_test_metrics)
print("Random Forest - test directional accuracy:", rf_dir_acc)

Random Forest - validation metrics: {'RMSE': np.float64(0.008777587376121151), 'MAE': 0.006000000000000001}
Random Forest - test metrics: {'RMSE': np.float64(0.014802500689598751), 'MAE': 0.009275967600000001}
Random Forest - test directional accuracy: 0.538034425319267

```

The Random Forest model has  $R^2 \approx -0.35$ . This is worse off than Linear Regression model.

```

# 10. ARIMA baseline (simplified, univariate)

# We'll model LogReturn as a univariate series
logret_series = df["LogReturn"].dropna()

# Align with df_sup range (so we're forecasting over the same overall period)
logret_aligned = logret_series.loc[df_sup.index]

train_end_idx = df_train.index[-1]
val_end_idx = df_val.index[-1]

```

```

logret_train = logret_aligned.loc[:train_end_idx]
logret_val = logret_aligned.loc[train_end_idx:val_end_idx]
logret_test = logret_aligned.loc[val_end_idx:]

# Fit a simple ARIMA(1,0,1) on the training portion
arima_model = ARIMA(logret_train, order=(1, 0, 1))
arima_result = arima_model.fit()

# Forecast over validation + test window
n_forecast = len(logret_val) + len(logret_test)
arima_forecast = arima_result.forecast(steps=n_forecast)

# Take the last n_forecast actual values to compare against
logret_val_test = logret_aligned.iloc[-n_forecast:]

# - - - Clean and align for metrics (avoid NaNs) - - -
# Put both into arrays
y_true = logret_val_test.to_numpy()
y_pred = np.asarray(arima_forecast)

# In case lengths differ for any reason, clip to the shorter one
min_len = min(len(y_true), len(y_pred))
y_true = y_true[:min_len]
y_pred = y_pred[:min_len]

# Drop any NaNs that might still be present
mask = (~np.isnan(y_true)) & (~np.isnan(y_pred))
y_true_clean = y_true[mask]
y_pred_clean = y_pred[mask]

# Compute metrics
arima_metrics = compute_regression_metrics(y_true_clean, y_pred_clean)
arima_dir_acc = directional_accuracy(y_true_clean, y_pred_clean)

print("ARIMA metrics (val+test window):", arima_metrics)
print("ARIMA directional accuracy:", arima_dir_acc)

# Plot actual vs forecast on the same time index
plot_index = logret_val_test.index[:len(y_true_clean)]
plt.plot(plot_index, y_true_clean, label="Actual")
plt.plot(plot_index, y_pred_clean, label="ARIMA forecast", alpha=0.7)
plt.legend()
plt.title("ARIMA Forecast vs Actual LogReturn")
plt.savefig("../images/ARIMA.png")
plt.show()


```

```

/home/jovyan/.local/share/envs/final -group13/lib/python3.10/site -packages/statsmodels/tsa/
    self._init_dates(dates, freq)
/home/jovyan/.local/share/envs/final -group13/lib/python3.10/site -packages/statsmodels/tsa/
    self._init_dates(dates, freq)
/home/jovyan/.local/share/envs/final -group13/lib/python3.10/site -packages/statsmodels/tsa/
    self._init_dates(dates, freq)
/home/jovyan/.local/share/envs/final -group13/lib/python3.10/site -packages/statsmodels/tsa/
    return get_prediction_index()
/home/jovyan/.local/share/envs/final -group13/lib/python3.10/site -packages/statsmodels/tsa/
    return get_prediction_index()

ARIMA metrics (val+test window): {'RMSE': np.float64(0.010885095016858535), 'MAE': 0.0071450
ARIMA directional accuracy: 0.5456563974465723

```



The ARIMA model has  $R^2 \approx -0.0009$ . This is far better than our linear baseline.

```

#LSTM Data Prep
tf.random.set_seed(42)
scaler = StandardScaler()
X_train_sc = scaler.fit_transform(X_train)
X_val_sc = scaler.transform(X_val)
X_test_sc = scaler.transform(X_test)

SEQ_LEN = 30
X_train_lstm, y_train_lstm = make_lstm_sequences(X_train_sc, y_train.values, SEQ_LEN)
X_val_lstm, y_val_lstm = make_lstm_sequences(X_val_sc, y_val.values, SEQ_LEN)
X_test_lstm, y_test_lstm = make_lstm_sequences(X_test_sc, y_test.values, SEQ_LEN)

#LSTM Build
model = models.Sequential([
    layers.Input(shape=(X_train_lstm.shape[1], X_train_lstm.shape[2])),
    layers.LSTM(32, return_sequences=False),
    layers.Dropout(0.2),
    layers.Dense(16, activation='relu'),
    layers.Dense(1)
]

```

```
] model.compile(optimizer='adam', loss='mse')

history = model.fit(
    X_train_lstm, y_train_lstm,
    validation_data=(X_val_lstm, y_val_lstm),
    epochs=15,
    batch_size=32,
    verbose=1
)

plt.figure()
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.title('LSTM Training Loss')
plt.legend()
plt.savefig("../images/lstm_training_curve.png")
plt.show()

2025-12-11 01:55:28.393109: E external/local_xla/xla/stream_executor/cuda/cuda_platform.cc:168 Epoch 1/15
168/168 4s 11ms/step - loss: 0.0066 - val_loss: 4.7753e-04
Epoch 2/15
168/168 2s 9ms/step - loss: 0.0012 - val_loss: 3.9340e-04
Epoch 3/15
168/168 2s 10ms/step - loss: 6.4524e-04 - val_loss: 1.7822e-04
Epoch 4/15
168/168 2s 10ms/step - loss: 4.0324e-04 - val_loss: 9.1113e-05
Epoch 5/15
168/168 2s 9ms/step - loss: 3.1936e-04 - val_loss: 7.6343e-05
Epoch 6/15
168/168 2s 9ms/step - loss: 2.5205e-04 - val_loss: 7.1253e-05
Epoch 7/15
168/168 2s 9ms/step - loss: 2.1402e-04 - val_loss: 6.9024e-05
Epoch 8/15
168/168 2s 9ms/step - loss: 2.0618e-04 - val_loss: 6.9460e-05
Epoch 9/15
168/168 2s 9ms/step - loss: 1.8815e-04 - val_loss: 7.0081e-05
Epoch 10/15
168/168 2s 9ms/step - loss: 1.7364e-04 - val_loss: 6.9981e-05
Epoch 11/15
168/168 2s 10ms/step - loss: 1.6579e-04 - val_loss: 6.9683e-05
Epoch 12/15
168/168 2s 9ms/step - loss: 1.5954e-04 - val_loss: 6.8659e-05
```

```

Epoch 13/15
168/168          2s 9ms/step - loss: 1.5293e -04 - val_loss: 6.8648e -05
Epoch 14/15
168/168          2s 9ms/step - loss: 1.5272e -04 - val_loss: 6.8502e -05
Epoch 15/15
168/168          2s 9ms/step - loss: 1.5327e -04 - val_loss: 6.8400e -05

```



```

#LSTM evaluations
y_pred_lstm = model.predict(X_test_lstm).flatten()

lstm_metrics = compute_regression_metrics(y_test_lstm, y_pred_lstm)
lstm_dir_acc = directional_accuracy(y_test_lstm, y_pred_lstm)

print("LSTM Metrics:", lstm_metrics)
print(f"Directional Accuracy: {lstm_dir_acc:.2%}")

56/56          0s 5ms/step
LSTM Metrics: {'RMSE': np.float64(0.01594281443908124), 'MAE': 0.010460818840851595, 'MAPE': 49.24%}

summary_rows = []

summary_rows.append(
    {"Model": "Naive",
     **baseline_metrics,
     "DirAcc": baseline_dir_acc}
)

summary_rows.append(
    {"Model": "LinearRegression",
     **lr_test_metrics,
     "DirAcc": lr_dir_acc}
)

summary_rows.append(

```

```

        {"Model": "RandomForest",
         **rf_test_metrics,
         "DirAcc": rf_dir_acc}
    )

summary_rows.append(
    {"Model": "ARIMA(1,0,1)",
     **arima_metrics,
     "DirAcc": arima_dir_acc}
)

summary_rows.append(
    {"Model": "LSTM",
     **lstm_metrics,
     "DirAcc": lstm_dir_acc}
)

summary_df = pd.DataFrame(summary_rows)
summary_df.to_csv("../data/final_metrics.csv")
summary_df.head()

```

	Model	RMSE	MAE	MAPE	R2	DirAcc
0	Naive	0.016941	0.011610	5.502865e+07	0.509162	0.766366
1	LinearReg	0.913780	0.010006	3.790434e+07	0.464187	0.168709
2	RandomFor	0.919296	0.015812	5.571611e+07	0.463631	1.291640
3	ARIMA(1,0,1)	0.010885	0.007145	3.433071e+07	0.545656	0.000894
4	LSTM	0.015943	0.010461	4.436545e+07	0.492377	0.573669

```

fig, axes = plt.subplots(1, 2, figsize=(12, 4))

axes[0].bar(summary_df["Model"], summary_df["R2"])
axes[0].set_title("R2 by Model")
axes[0].set_xticklabels(summary_df["Model"], rotation=45, ha="right")

axes[1].bar(summary_df["Model"], summary_df["DirAcc"])
axes[1].set_title("Directional Accuracy by Model")
axes[1].set_xticklabels(summary_df["Model"], rotation=45, ha="right")

plt.tight_layout()


```

```

plt.savefig("../images/R2_Accuracy.png")
plt.show()

/tmp/ipykernel_753/602907317.py:5: UserWarning: set_ticklabels() should only be used with a
  axes[0].set_xticklabels(summary_df["Model"], rotation=45, ha="right")
/tmp/ipykernel_753/602907317.py:9: UserWarning: set_ticklabels() should only be used with a
  axes[1].set_xticklabels(summary_df["Model"], rotation=45, ha="right")

```



All our models find it difficult to capture the true nature of returns in the stock market. However, ARIMA clearly is the best model.