

# Model Training and Evaluating it

Wednesday 10<sup>th</sup> December, 2025



**Description:** This notebook creates various models and evaluates them as needed.

```
#Setting everything up
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import os
import sys

from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.preprocessing import StandardScaler
from statsmodels.tsa.arima.model import ARIMA
import tensorflow as tf
from tensorflow.keras import layers, models

sys.path.append(os.path.abspath(os.path.join('..')))

from scripts.data_process import make_supervised_frame, time_series_split
from scripts.data_modeling import compute_regression_metrics, directional_accuracy, make_lstm

np.random.seed(42)
tf.random.set_seed(42)
plt.rcParams["figure.figsize"] = (10, 5)
plt.rcParams["axes.grid"] = True

2025 -12 -10 01:00:04.138126: I external/local_xla/xla/tsl/cuda/cudart_stub.cc:31] Could not
2025 -12 -10 01:00:04.138576: I tensorflow/core/util/port.cc:153] oneDNN custom operations a
2025 -12 -10 01:00:04.199951: I tensorflow/core/platform/cpu_feature_guard.cc:210] This Tens
To enable the following instructions: AVX2 AVX512F AVX512_VNNI FMA, in other operations, reb

2025 -12 -10 01:00:06.801074: I tensorflow/core/util/port.cc:153] oneDNN custom operations a
2025 -12 -10 01:00:06.802886: I external/local_xla/xla/tsl/cuda/cudart_stub.cc:31] Could not
```

```
#Loading data
df = pd.read_csv("../data/sp500.csv", index_col = "Date", parse_dates = True)

df_sup = make_supervised_frame(df, target_col="LogReturn", horizon=1, lags=3)

df_sup.head()
```

	Clos	Hig	Low	Ope	Volum	MA1	MA5	MA10	MA50	LogRet	Ret_lag1	Ret_lag2	Ret_lag3	Date
1990-03-13	336.880	337.035	335.917	336.030	885.335	335.923	336.140	0.005620	0.922821	0.002582	0.187	0.009805		
								0.007894	7915		0.212164		0.006901	
1990-03-14	336.887	337.030	336.916	336.045	900.000	335.923	336.631	0.005610	0.922760	0.002584	0.187	0.009805		
										0.007915		0.006901		
1990-03-15	338.878	337.035	336.915	337.015	907.000	337.030	338.631	0.005617	0.922763	0.002584	0.187	0.009805		
										0.007915		0.006901		
1990-03-16	341.940	938.037	340.915	342.022	927.000	340.934	341.631	0.005617	0.922763	0.002584	0.187	0.009805		
										0.007915		0.006901		
1990-03-19	343.329	939.039	340.912	340.808	903.000	346.237	347.301	0.005617	0.922767	0.002584	0.187	0.009805		
										0.005722		0.006901		

```
df_train, df_val, df_test = time_series_split(df_sup, train_frac=0.6, val_frac=0.2)

feature_cols = [c for c in df_sup.columns if c !=["y"]]

X_train, y_train = df_train[feature_cols], df_train["y"]
X_val, y_val = df_val[feature_cols], df_val["y"]
X_test, y_test = df_test[feature_cols], df_test["y"]

print("Train size:", len(X_train))
print("Val size:", len(X_val))
print("Test size:", len(X_test))

Train size: 5401
Val size: 1800
Test size: 1801

#baseline
y_test_naive = df_test["LogReturn_lag1"]
```

```

baseline_metrics = compute_regression_metrics(y_test, y_test_naive)
baseline_dir_acc = directional_accuracy(y_test, y_test_naive)

print("Naive baseline metrics:", baseline_metrics)
print("Naive baseline directional accuracy:", baseline_dir_acc)

Naive baseline metrics: {'RMSE': np.float64(0.016940516181126818), 'MAE': 0.0116095086299034}
Naive baseline directional accuracy: 0.5091615769017213

# Linear Regression
linreg = LinearRegression()
linreg.fit(X_train, y_train)

y_val_lr = linreg.predict(X_val)
y_test_lr = linreg.predict(X_test)

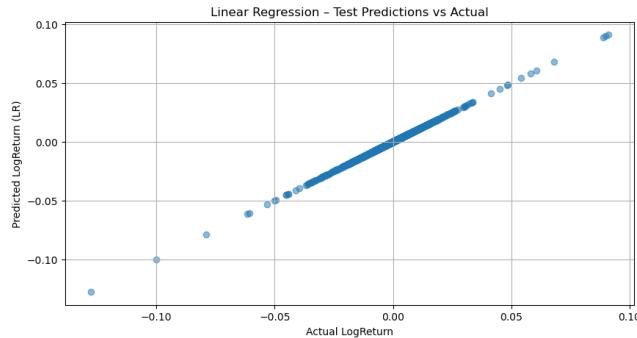
lr_val_metrics = compute_regression_metrics(y_val, y_val_lr)
lr_test_metrics = compute_regression_metrics(y_test, y_test_lr)
lr_dir_acc = directional_accuracy(y_test, y_test_lr)

print("Linear Regression - Validation Metrics:", lr_val_metrics)
print("Linear Regression - Test Metrics:", lr_test_metrics)
print(f"Linear Regression - Directional Accuracy: {lr_dir_acc:.2%}")

Linear Regression - Validation Metrics: {'RMSE': np.float64(7.178125137080271e-08), 'MAE': 1.411111111111111}
Linear Regression - Test Metrics: {'RMSE': np.float64(1.9145641052797167e-07), 'MAE': 1.411111111111111}
Linear Regression - Directional Accuracy: 100.00%

plt.scatter(y_test, y_test_lr, alpha=0.5)
plt.xlabel("Actual LogReturn")
plt.ylabel("Predicted LogReturn (LR)")
plt.title("Linear Regression - Test Predictions vs Actual")
plt.savefig("../images/LR.png")
plt.show()

```



```

rf = RandomForestRegressor(n_estimators=200, max_depth=None, random_state=42)
rf.fit(X_train, y_train)

y_val_rf = rf.predict(X_val)
y_test_rf = rf.predict(X_test)

rf_val_metrics = compute_regression_metrics(y_val, y_val_rf)
rf_test_metrics = compute_regression_metrics(y_test, y_test_rf)
rf_dir_acc = directional_accuracy(y_test, y_test_rf)

print("Random Forest - validation metrics:", rf_val_metrics)
print("Random Forest - test metrics:", rf_test_metrics)
print("Random Forest - test directional accuracy:", rf_dir_acc)

Random Forest - validation metrics: {'RMSE': np.float64(3.989333611068483e -05), 'MAE': 8.73}
Random Forest - test metrics: {'RMSE': np.float64(0.0011719469383807014), 'MAE': 8.215782319}
Random Forest - test directional accuracy: 0.9994447529150472

# 10. ARIMA baseline (simplified, univariate)

# We'll model LogReturn as a univariate series
logret_series = df["LogReturn"].dropna()

# Align with df_sup range (so we're forecasting over the same overall period)
logret_aligned = logret_series.loc[df_sup.index]

train_end_idx = df_train.index[ -1]
val_end_idx = df_val.index[ -1]

logret_train = logret_aligned.loc[:train_end_idx]
logret_val = logret_aligned.loc[train_end_idx:val_end_idx]
logret_test = logret_aligned.loc[val_end_idx:]

# Fit a simple ARIMA(1,0,1) on the training portion
arima_model = ARIMA(logret_train, order=(1, 0, 1))
arima_result = arima_model.fit()

# Forecast over validation + test window
n_forecast = len(logret_val) + len(logret_test)
arima_forecast = arima_result.forecast(steps=n_forecast)

# Take the last n_forecast actual values to compare against
logret_val_test = logret_aligned.iloc[ -n_forecast:]

# - - - Clean and align for metrics (avoid NaNs) - - -

```

```

# Put both into arrays
y_true = logret_val_test.to_numpy()
y_pred = np.asarray(arima_forecast)

# In case lengths differ for any reason, clip to the shorter one
min_len = min(len(y_true), len(y_pred))
y_true = y_true[:min_len]
y_pred = y_pred[:min_len]

# Drop any NaNs that might still be present
mask = (~np.isnan(y_true)) & (~np.isnan(y_pred))
y_true_clean = y_true[mask]
y_pred_clean = y_pred[mask]

# Compute metrics
arima_metrics = compute_regression_metrics(y_true_clean, y_pred_clean)
arima_dir_acc = directional_accuracy(y_true_clean, y_pred_clean)

print("ARIMA metrics (val+test window):", arima_metrics)
print("ARIMA directional accuracy:", arima_dir_acc)

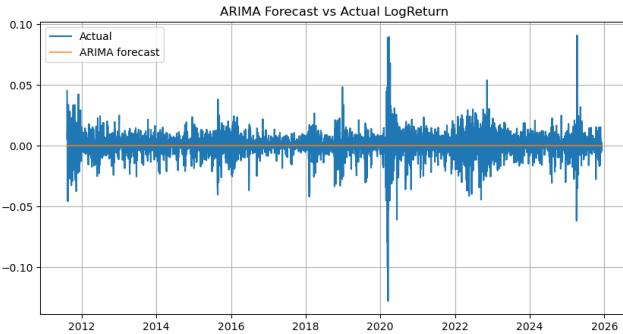
# Plot actual vs forecast on the same time index
plot_index = logret_val_test.index[:len(y_true_clean)]
plt.plot(plot_index, y_true_clean, label="Actual")
plt.plot(plot_index, y_pred_clean, label="ARIMA forecast", alpha=0.7)
plt.legend()
plt.title("ARIMA Forecast vs Actual LogReturn")
plt.savefig("../images/ARIMA.png")
plt.show()

/home/jovyan/.local/share/envs/final -group13/lib/python3.10/site -packages/statsmodels/tsa
    self._init_dates(dates, freq)
/home/jovyan/.local/share/envs/final -group13/lib/python3.10/site -packages/statsmodels/tsa
    self._init_dates(dates, freq)
/home/jovyan/.local/share/envs/final -group13/lib/python3.10/site -packages/statsmodels/tsa
    self._init_dates(dates, freq)

/home/jovyan/.local/share/envs/final -group13/lib/python3.10/site -packages/statsmodels/tsa
    return get_prediction_index(
/home/jovyan/.local/share/envs/final -group13/lib/python3.10/site -packages/statsmodels/tsa
    return get_prediction_index(


ARIMA metrics (val+test window): {'RMSE': np.float64(0.010885095016858535), 'MAE': 0.0071450
ARIMA directional accuracy: 0.5456563974465723

```



```

#LSTM Data Prep
tf.random.set_seed(42)
scaler = StandardScaler()
X_train_sc = scaler.fit_transform(X_train)
X_val_sc = scaler.transform(X_val)
X_test_sc = scaler.transform(X_test)

SEQ_LEN = 30
X_train_lstm, y_train_lstm = make_lstm_sequences(X_train_sc, y_train.values, SEQ_LEN)
X_val_lstm, y_val_lstm = make_lstm_sequences(X_val_sc, y_val.values, SEQ_LEN)
X_test_lstm, y_test_lstm = make_lstm_sequences(X_test_sc, y_test.values, SEQ_LEN)

#LSTM Build
model = models.Sequential([
    layers.Input(shape=(X_train_lstm.shape[1], X_train_lstm.shape[2])),
    layers.LSTM(32, return_sequences=False),
    layers.Dropout(0.2),
    layers.Dense(16, activation='relu'),
    layers.Dense(1)
])

model.compile(optimizer='adam', loss='mse')

history = model.fit(
    X_train_lstm, y_train_lstm,
    validation_data=(X_val_lstm, y_val_lstm),
    epochs=15,
    batch_size=32,
    verbose=1
)

plt.figure()
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.title('LSTM Training Loss')

```

```
plt.legend()
plt.savefig("../images/lstm_training_curve.png")
plt.show()

Epoch 1/15

2025 -12 -10 01:00:27.758075: E external/local_xla/xla/stream_executor/cuda/cuda_platform.cc:168:168       4:18 2s/step - loss: 0.0659

6/168           1s 10ms/step - loss: 0.0519

12/168          1s 10ms/step - loss: 0.0434

18/168          1s 10ms/step - loss: 0.0383

24/168          1s 9ms/step - loss: 0.0349

30/168          1s 9ms/step - loss: 0.0322

36/168          1s 9ms/step - loss: 0.0301

42/168          1s 9ms/step - loss: 0.0284

48/168          1s 9ms/step - loss: 0.0269

54/168          1s 9ms/step - loss: 0.0257

60/168          0s 9ms/step - loss: 0.0246

66/168          0s 9ms/step - loss: 0.0236

72/168          0s 9ms/step - loss: 0.0227

78/168          0s 9ms/step - loss: 0.0219
```

84/168            0s 9ms/step - loss: 0.0212

90/168            0s 9ms/step - loss: 0.0205

96/168            0s 9ms/step - loss: 0.0199

102/168           0s 9ms/step - loss: 0.0193

108/168           0s 9ms/step - loss: 0.0188

114/168           0s 9ms/step - loss: 0.0183

120/168           0s 9ms/step - loss: 0.0179

126/168           0s 9ms/step - loss: 0.0174

132/168           0s 9ms/step - loss: 0.0170

138/168           0s 9ms/step - loss: 0.0166

144/168           0s 9ms/step - loss: 0.0163

150/168           0s 9ms/step - loss: 0.0159

157/168           0s 9ms/step - loss: 0.0156

163/168           0s 9ms/step - loss: 0.0153

168/168           4s 12ms/step - loss: 0.0071 - val\_loss: 2.7630e -04

Epoch 2/15

1/168            4s 28ms/step - loss: 0.0015

7/168            1s 9ms/step - loss: 0.0020

13/168            1s 9ms/step - loss: 0.0020

19/168            1s 9ms/step - loss: 0.0020

25/168            1s 9ms/step - loss: 0.0019

31/168            1s 9ms/step - loss: 0.0019

37/168            1s 9ms/step - loss: 0.0018

43/168            1s 9ms/step - loss: 0.0018

49/168            1s 9ms/step - loss: 0.0018

55/168            1s 9ms/step - loss: 0.0018

61/168            0s 9ms/step - loss: 0.0018

67/168            0s 9ms/step - loss: 0.0017

73/168            0s 9ms/step - loss: 0.0017

79/168            0s 9ms/step - loss: 0.0017

85/168            0s 9ms/step - loss: 0.0017

91/168            0s 9ms/step - loss: 0.0017

97/168            0s 9ms/step - loss: 0.0017

103/168           0s 9ms/step - loss: 0.0016

109/168            0s 9ms/step - loss: 0.0016

115/168            0s 9ms/step - loss: 0.0016

121/168            0s 9ms/step - loss: 0.0016

127/168            0s 9ms/step - loss: 0.0016

133/168            0s 9ms/step - loss: 0.0016

139/168            0s 9ms/step - loss: 0.0016

145/168            0s 9ms/step - loss: 0.0015

151/168            0s 9ms/step - loss: 0.0015

157/168            0s 9ms/step - loss: 0.0015

163/168            0s 9ms/step - loss: 0.0015

168/168            2s 10ms/step - loss: 0.0011 - val\_loss: 1.3414e -04

Epoch 3/15

1/168            5s 30ms/step - loss: 6.7609e -04

7/168            1s 9ms/step - loss: 7.5495e -04

13/168            1s 9ms/step - loss: 7.2350e -04

19/168            1s 9ms/step - loss: 6.9969e -04

25/168            1s 9ms/step - loss: 6.9050e -04

31/168            1s 9ms/step - loss: 6.8531e -04

37/168            1s 9ms/step - loss: 6.8148e -04

43/168            1s 9ms/step - loss: 6.7856e -04

49/168            1s 9ms/step - loss: 6.7763e -04

55/168            1s 9ms/step - loss: 6.7576e -04

61/168            0s 9ms/step - loss: 6.7458e -04

67/168            0s 9ms/step - loss: 6.7229e -04

73/168            0s 9ms/step - loss: 6.7041e -04

79/168            0s 9ms/step - loss: 6.6839e -04

85/168            0s 9ms/step - loss: 6.6559e -04

91/168            0s 9ms/step - loss: 6.6239e -04

97/168            0s 9ms/step - loss: 6.5863e -04

103/168           0s 9ms/step - loss: 6.5498e -04

109/168           0s 9ms/step - loss: 6.5124e -04

115/168           0s 9ms/step - loss: 6.4737e -04

121/168           0s 9ms/step - loss: 6.4362e -04

127/168           0s 9ms/step - loss: 6.3986e -04

133/168            0s 9ms/step - loss: 6.3634e -04

139/168            0s 9ms/step - loss: 6.3279e -04

145/168            0s 9ms/step - loss: 6.2934e -04

151/168            0s 9ms/step - loss: 6.2599e -04

157/168            0s 9ms/step - loss: 6.2266e -04

163/168            0s 9ms/step - loss: 6.1934e -04

168/168            2s 10ms/step - loss: 5.2956e -04 - val\_loss: 8.0141e -05

Epoch 4/15

1/168            5s 32ms/step - loss: 4.4027e -04

7/168            1s 9ms/step - loss: 4.1018e -04

13/168            1s 9ms/step - loss: 3.9620e -04

19/168            1s 9ms/step - loss: 3.9031e -04

25/168            1s 9ms/step - loss: 3.9210e -04

31/168            1s 9ms/step - loss: 3.9565e -04

37/168            1s 9ms/step - loss: 3.9872e -04

43/168            1s 9ms/step - loss: 4.0290e -04

49/168            1s 9ms/step - loss: 4.0694e -04

55/168            1s 9ms/step - loss: 4.1017e -04

61/168            0s 9ms/step - loss: 4.1593e -04

67/168            0s 9ms/step - loss: 4.1909e -04

73/168            0s 9ms/step - loss: 4.2113e -04

79/168            0s 9ms/step - loss: 4.2290e -04

85/168            0s 9ms/step - loss: 4.2409e -04

91/168            0s 9ms/step - loss: 4.2467e -04

97/168            0s 9ms/step - loss: 4.2449e -04

103/168           0s 9ms/step - loss: 4.2414e -04

109/168           0s 9ms/step - loss: 4.2380e -04

115/168           0s 9ms/step - loss: 4.2334e -04

121/168           0s 9ms/step - loss: 4.2278e -04

127/168           0s 9ms/step - loss: 4.2203e -04

133/168           0s 9ms/step - loss: 4.2110e -04

139/168           0s 9ms/step - loss: 4.2002e -04

145/168           0s 9ms/step - loss: 4.1889e -04

151/168           0s 9ms/step - loss: 4.1767e -04

157/168            0s 9ms/step - loss: 4.1646e -04

163/168            0s 9ms/step - loss: 4.1521e -04

168/168            2s 11ms/step - loss: 3.7602e -04 - val\_loss: 6.9167e -05

Epoch 5/15

1/168            5s 30ms/step - loss: 4.9068e -04

7/168            1s 9ms/step - loss: 3.1018e -04

13/168            1s 9ms/step - loss: 2.9280e -04

19/168            1s 9ms/step - loss: 2.9199e -04

25/168            1s 9ms/step - loss: 2.8865e -04

31/168            1s 9ms/step - loss: 2.8467e -04

37/168            1s 9ms/step - loss: 2.8300e -04

43/168            1s 9ms/step - loss: 2.8519e -04

49/168            1s 9ms/step - loss: 2.8850e -04

55/168            1s 9ms/step - loss: 2.9163e -04

61/168            0s 9ms/step - loss: 2.9433e -04

67/168            0s 9ms/step - loss: 2.9620e -04

73/168            0s 9ms/step - loss: 2.9711e -04

79/168            0s 9ms/step - loss: 2.9758e -04

85/168            0s 9ms/step - loss: 2.9776e -04

91/168            0s 9ms/step - loss: 2.9754e -04

97/168            0s 9ms/step - loss: 2.9693e -04

103/168           0s 9ms/step - loss: 2.9628e -04

109/168           0s 9ms/step - loss: 2.9567e -04

115/168           0s 9ms/step - loss: 2.9505e -04

121/168           0s 9ms/step - loss: 2.9451e -04

127/168           0s 9ms/step - loss: 2.9395e -04

133/168           0s 9ms/step - loss: 2.9327e -04

139/168           0s 9ms/step - loss: 2.9265e -04

144/168           0s 9ms/step - loss: 2.9217e -04

150/168           0s 9ms/step - loss: 2.9152e -04

156/168           0s 9ms/step - loss: 2.9088e -04

162/168           0s 9ms/step - loss: 2.9026e -04

168/168           0s 9ms/step - loss: 2.8965e -04

168/168           2s 10ms/step - loss: 2.7244e -04 - val\_loss: 6.8015e -05

Epoch 6/15

1/168	5s 30ms/step - loss: 2.5459e -04
7/168	1s 9ms/step - loss: 2.3608e -04
13/168	1s 9ms/step - loss: 2.2994e -04
19/168	1s 9ms/step - loss: 2.3191e -04
25/168	1s 9ms/step - loss: 2.3113e -04
31/168	1s 9ms/step - loss: 2.2987e -04
37/168	1s 9ms/step - loss: 2.2966e -04
43/168	1s 9ms/step - loss: 2.2930e -04
49/168	1s 9ms/step - loss: 2.2900e -04
55/168	1s 9ms/step - loss: 2.2934e -04
61/168	0s 9ms/step - loss: 2.2971e -04
67/168	0s 9ms/step - loss: 2.2983e -04
73/168	0s 9ms/step - loss: 2.2962e -04
79/168	0s 9ms/step - loss: 2.2940e -04
85/168	0s 9ms/step - loss: 2.2920e -04
91/168	0s 9ms/step - loss: 2.2880e -04

97/168            0s 9ms/step - loss: 2.2824e -04

103/168            0s 9ms/step - loss: 2.2771e -04

109/168            0s 9ms/step - loss: 2.2734e -04

115/168            0s 9ms/step - loss: 2.2711e -04

121/168            0s 9ms/step - loss: 2.2699e -04

127/168            0s 9ms/step - loss: 2.2677e -04

133/168            0s 9ms/step - loss: 2.2652e -04

139/168            0s 9ms/step - loss: 2.2623e -04

145/168            0s 9ms/step - loss: 2.2600e -04

151/168            0s 9ms/step - loss: 2.2578e -04

157/168            0s 9ms/step - loss: 2.2555e -04

163/168            0s 9ms/step - loss: 2.2537e -04

168/168            2s 10ms/step - loss: 2.2120e -04 - val\_loss: 7.0077e -05

Epoch 7/15

1/168            5s 33ms/step - loss: 2.9188e -04

6/168            1s 10ms/step - loss: 2.3806e -04

12/168            1s 9ms/step - loss: 2.3674e -04

18/168	1s 9ms/step - loss: 2.3599e -04
24/168	1s 9ms/step - loss: 2.3220e -04
30/168	1s 9ms/step - loss: 2.2813e -04
36/168	1s 9ms/step - loss: 2.2600e -04
42/168	1s 9ms/step - loss: 2.2464e -04
48/168	1s 9ms/step - loss: 2.2352e -04
54/168	1s 9ms/step - loss: 2.2263e -04
60/168	0s 9ms/step - loss: 2.2212e -04
66/168	0s 9ms/step - loss: 2.2106e -04
72/168	0s 9ms/step - loss: 2.1992e -04
78/168	0s 9ms/step - loss: 2.1896e -04
84/168	0s 9ms/step - loss: 2.1803e -04
90/168	0s 9ms/step - loss: 2.1706e -04
96/168	0s 9ms/step - loss: 2.1600e -04
102/168	0s 9ms/step - loss: 2.1516e -04
108/168	0s 9ms/step - loss: 2.1454e -04
114/168	0s 9ms/step - loss: 2.1403e -04

120/168            0s 9ms/step - loss: 2.1367e -04

126/168            0s 9ms/step - loss: 2.1335e -04

132/168            0s 9ms/step - loss: 2.1301e -04

137/168            0s 9ms/step - loss: 2.1270e -04

143/168            0s 9ms/step - loss: 2.1238e -04

149/168            0s 9ms/step - loss: 2.1204e -04

155/168            0s 9ms/step - loss: 2.1168e -04

161/168            0s 9ms/step - loss: 2.1134e -04

167/168            0s 9ms/step - loss: 2.1101e -04

168/168            2s 11ms/step - loss: 2.0142e -04 - val\_loss: 7.3925e -05

Epoch 8/15

1/168            5s 30ms/step - loss: 2.5238e -04

7/168            1s 10ms/step - loss: 1.8993e -04

13/168            1s 9ms/step - loss: 1.8659e -04

19/168            1s 9ms/step - loss: 1.8877e -04

25/168            1s 9ms/step - loss: 1.8877e -04

31/168            1s 9ms/step - loss: 1.8923e -04

37/168            1s 9ms/step - loss: 1.9008e -04

43/168            1s 9ms/step - loss: 1.9012e -04

49/168            1s 9ms/step - loss: 1.9001e -04

55/168            1s 9ms/step - loss: 1.9023e -04

61/168            0s 9ms/step - loss: 1.9125e -04

67/168            0s 9ms/step - loss: 1.9170e -04

73/168            0s 9ms/step - loss: 1.9170e -04

79/168            0s 9ms/step - loss: 1.9157e -04

85/168            0s 9ms/step - loss: 1.9143e -04

91/168            0s 9ms/step - loss: 1.9112e -04

97/168            0s 9ms/step - loss: 1.9069e -04

103/168           0s 9ms/step - loss: 1.9035e -04

109/168           0s 9ms/step - loss: 1.9016e -04

115/168           0s 9ms/step - loss: 1.9009e -04

121/168           0s 9ms/step - loss: 1.9014e -04

127/168           0s 9ms/step - loss: 1.9018e -04

133/168           0s 9ms/step - loss: 1.9023e -04

139/168            0s 9ms/step - loss: 1.9022e -04

145/168            0s 9ms/step - loss: 1.9021e -04

151/168            0s 9ms/step - loss: 1.9020e -04

157/168            0s 9ms/step - loss: 1.9019e -04

163/168            0s 9ms/step - loss: 1.9015e -04

168/168            2s 10ms/step - loss: 1.8843e -04 - val\_loss: 7.0693e -05

Epoch 9/15

1/168            5s 30ms/step - loss: 3.2748e -04

7/168            1s 9ms/step - loss: 2.2604e -04

13/168            1s 9ms/step - loss: 2.0402e -04

19/168            1s 9ms/step - loss: 1.9686e -04

25/168            1s 9ms/step - loss: 1.9264e -04

31/168            1s 9ms/step - loss: 1.8949e -04

37/168            1s 9ms/step - loss: 1.8783e -04

43/168            1s 9ms/step - loss: 1.8631e -04

49/168            1s 9ms/step - loss: 1.8562e -04

55/168            1s 9ms/step - loss: 1.8567e -04

61/168	0s 9ms/step - loss: 1.8569e -04
67/168	0s 9ms/step - loss: 1.8510e -04
73/168	0s 9ms/step - loss: 1.8433e -04
78/168	0s 9ms/step - loss: 1.8369e -04
84/168	0s 9ms/step - loss: 1.8290e -04
90/168	0s 9ms/step - loss: 1.8207e -04
96/168	0s 9ms/step - loss: 1.8128e -04
102/168	0s 9ms/step - loss: 1.8059e -04
108/168	0s 9ms/step - loss: 1.8011e -04
114/168	0s 9ms/step - loss: 1.7975e -04
120/168	0s 9ms/step - loss: 1.7961e -04
126/168	0s 9ms/step - loss: 1.7949e -04
132/168	0s 9ms/step - loss: 1.7934e -04
138/168	0s 9ms/step - loss: 1.7917e -04
144/168	0s 9ms/step - loss: 1.7904e -04
150/168	0s 9ms/step - loss: 1.7888e -04
156/168	0s 9ms/step - loss: 1.7870e -04

162/168            0s 9ms/step - loss: 1.7857e -04

168/168            0s 9ms/step - loss: 1.7850e -04

168/168            2s 11ms/step - loss: 1.7658e -04 - val\_loss: 6.9025e -05

Epoch 10/15

1/168            5s 31ms/step - loss: 2.5594e -04

7/168            1s 9ms/step - loss: 1.8167e -04

13/168            1s 9ms/step - loss: 1.7392e -04

19/168            1s 9ms/step - loss: 1.7335e -04

25/168            1s 9ms/step - loss: 1.7149e -04

31/168            1s 9ms/step - loss: 1.7062e -04

37/168            1s 9ms/step - loss: 1.7073e -04

43/168            1s 9ms/step - loss: 1.7076e -04

49/168            1s 9ms/step - loss: 1.7074e -04

55/168            1s 9ms/step - loss: 1.7107e -04

61/168            0s 9ms/step - loss: 1.7178e -04

67/168            0s 9ms/step - loss: 1.7201e -04

73/168            0s 9ms/step - loss: 1.7191e -04

79/168            0s 9ms/step - loss: 1.7172e -04

85/168            0s 9ms/step - loss: 1.7146e -04

91/168            0s 9ms/step - loss: 1.7104e -04

97/168            0s 9ms/step - loss: 1.7057e -04

103/168           0s 9ms/step - loss: 1.7022e -04

109/168           0s 9ms/step - loss: 1.7007e -04

115/168           0s 9ms/step - loss: 1.7015e -04

121/168           0s 9ms/step - loss: 1.7036e -04

127/168           0s 9ms/step - loss: 1.7056e -04

134/168           0s 9ms/step - loss: 1.7067e -04

141/168           0s 9ms/step - loss: 1.7073e -04

147/168           0s 9ms/step - loss: 1.7082e -04

153/168           0s 9ms/step - loss: 1.7087e -04

159/168           0s 9ms/step - loss: 1.7094e -04

165/168           0s 9ms/step - loss: 1.7098e -04

168/168           2s 10ms/step - loss: 1.7079e -04 - val\_loss: 6.7987e -05

Epoch 11/15

1/168            5s 30ms/step - loss: 3.4878e -04

7/168            1s 9ms/step - loss: 2.1028e -04

13/168           1s 9ms/step - loss: 1.8907e -04

19/168           1s 9ms/step - loss: 1.8030e -04

25/168           1s 9ms/step - loss: 1.7460e -04

31/168           1s 9ms/step - loss: 1.7076e -04

37/168           1s 9ms/step - loss: 1.6925e -04

43/168           1s 9ms/step - loss: 1.6776e -04

49/168           1s 9ms/step - loss: 1.6662e -04

55/168           1s 9ms/step - loss: 1.6605e -04

61/168           0s 9ms/step - loss: 1.6613e -04

67/168           0s 9ms/step - loss: 1.6595e -04

73/168           0s 9ms/step - loss: 1.6565e -04

79/168           0s 9ms/step - loss: 1.6531e -04

85/168           0s 9ms/step - loss: 1.6494e -04

91/168           0s 9ms/step - loss: 1.6446e -04

97/168           0s 9ms/step - loss: 1.6397e -04

103/168            0s 9ms/step - loss: 1.6365e -04

109/168            0s 9ms/step - loss: 1.6353e -04

115/168            0s 9ms/step - loss: 1.6359e -04

121/168            0s 9ms/step - loss: 1.6382e -04

127/168            0s 9ms/step - loss: 1.6396e -04

133/168            0s 9ms/step - loss: 1.6404e -04

139/168            0s 9ms/step - loss: 1.6407e -04

145/168            0s 9ms/step - loss: 1.6411e -04

151/168            0s 9ms/step - loss: 1.6411e -04

157/168            0s 9ms/step - loss: 1.6410e -04

163/168            0s 9ms/step - loss: 1.6407e -04

168/168            2s 10ms/step - loss: 1.6282e -04 - val\_loss: 6.9519e -05

Epoch 12/15

1/168            5s 33ms/step - loss: 1.8450e -04

7/168            1s 9ms/step - loss: 1.4972e -04

13/168            1s 9ms/step - loss: 1.4701e -04

19/168            1s 9ms/step - loss: 1.4861e -04

25/168            1s 9ms/step - loss: 1.4878e -04

31/168            1s 9ms/step - loss: 1.4870e -04

37/168            1s 9ms/step - loss: 1.4944e -04

43/168            1s 9ms/step - loss: 1.4994e -04

49/168            1s 9ms/step - loss: 1.5040e -04

55/168            0s 9ms/step - loss: 1.5123e -04

61/168            0s 9ms/step - loss: 1.5264e -04

67/168            0s 9ms/step - loss: 1.5336e -04

73/168            0s 9ms/step - loss: 1.5372e -04

79/168            0s 9ms/step - loss: 1.5392e -04

85/168            0s 9ms/step - loss: 1.5398e -04

91/168            0s 9ms/step - loss: 1.5390e -04

97/168            0s 9ms/step - loss: 1.5377e -04

103/168           0s 9ms/step - loss: 1.5375e -04

109/168           0s 9ms/step - loss: 1.5391e -04

115/168           0s 9ms/step - loss: 1.5422e -04

121/168           0s 9ms/step - loss: 1.5464e -04

127/168            0s 9ms/step - loss: 1.5502e -04

132/168            0s 9ms/step - loss: 1.5525e -04

138/168            0s 9ms/step - loss: 1.5549e -04

144/168            0s 9ms/step - loss: 1.5575e -04

150/168            0s 9ms/step - loss: 1.5599e -04

156/168            0s 9ms/step - loss: 1.5621e -04

162/168            0s 9ms/step - loss: 1.5642e -04

168/168            0s 9ms/step - loss: 1.5660e -04

168/168            2s 10ms/step - loss: 1.6101e -04 - val\_loss: 6.8521e -05

Epoch 13/15

1/168            4s 28ms/step - loss: 2.2210e -04

7/168            1s 9ms/step - loss: 1.6654e -04

13/168            1s 9ms/step - loss: 1.6005e -04

19/168            1s 9ms/step - loss: 1.5796e -04

25/168            1s 9ms/step - loss: 1.5574e -04

31/168            1s 9ms/step - loss: 1.5437e -04

37/168            1s 9ms/step - loss: 1.5426e -04

43/168            1s 9ms/step - loss: 1.5402e -04

49/168            1s 9ms/step - loss: 1.5375e -04

55/168            0s 9ms/step - loss: 1.5396e -04

62/168            0s 9ms/step - loss: 1.5450e -04

68/168            0s 9ms/step - loss: 1.5447e -04

74/168            0s 9ms/step - loss: 1.5432e -04

80/168            0s 9ms/step - loss: 1.5421e -04

86/168            0s 9ms/step - loss: 1.5397e -04

92/168            0s 9ms/step - loss: 1.5371e -04

98/168            0s 9ms/step - loss: 1.5345e -04

104/168           0s 9ms/step - loss: 1.5328e -04

110/168           0s 9ms/step - loss: 1.5325e -04

116/168           0s 9ms/step - loss: 1.5343e -04

122/168           0s 9ms/step - loss: 1.5369e -04

128/168           0s 9ms/step - loss: 1.5386e -04

134/168           0s 9ms/step - loss: 1.5394e -04

140/168           0s 9ms/step - loss: 1.5398e -04

146/168            0s 9ms/step - loss: 1.5403e -04

152/168            0s 9ms/step - loss: 1.5405e -04

158/168            0s 9ms/step - loss: 1.5408e -04

164/168            0s 9ms/step - loss: 1.5410e -04

168/168            2s 10ms/step - loss: 1.5406e -04 - val\_loss: 6.7993e -05

Epoch 14/15

1/168            5s 30ms/step - loss: 2.2555e -04

7/168            1s 9ms/step - loss: 1.6614e -04

13/168            1s 9ms/step - loss: 1.5605e -04

19/168            1s 9ms/step - loss: 1.5096e -04

25/168            1s 9ms/step - loss: 1.4884e -04

31/168            1s 9ms/step - loss: 1.4801e -04

37/168            1s 9ms/step - loss: 1.4815e -04

43/168            1s 9ms/step - loss: 1.4804e -04

49/168            1s 9ms/step - loss: 1.4807e -04

55/168            0s 9ms/step - loss: 1.4855e -04

61/168            0s 9ms/step - loss: 1.4924e -04

67/168            0s 9ms/step - loss: 1.4950e -04

73/168            0s 9ms/step - loss: 1.4962e -04

79/168            0s 9ms/step - loss: 1.4968e -04

85/168            0s 9ms/step - loss: 1.4964e -04

91/168            0s 9ms/step - loss: 1.4945e -04

97/168            0s 9ms/step - loss: 1.4928e -04

103/168           0s 9ms/step - loss: 1.4920e -04

109/168           0s 9ms/step - loss: 1.4931e -04

115/168           0s 9ms/step - loss: 1.4955e -04

121/168           0s 9ms/step - loss: 1.4992e -04

127/168           0s 9ms/step - loss: 1.5022e -04

133/168           0s 9ms/step - loss: 1.5047e -04

139/168           0s 9ms/step - loss: 1.5068e -04

145/168           0s 9ms/step - loss: 1.5090e -04

151/168           0s 9ms/step - loss: 1.5107e -04

157/168           0s 9ms/step - loss: 1.5121e -04

163/168           0s 9ms/step - loss: 1.5136e -04

168/168            2s 10ms/step - loss: 1.5460e -04 - val\_loss: 6.8454e -05  
Epoch 15/15  
1/168            5s 31ms/step - loss: 1.8362e -04  
  
7/168            1s 9ms/step - loss: 1.4532e -04  
  
13/168            1s 9ms/step - loss: 1.4250e -04  
  
19/168            1s 9ms/step - loss: 1.4203e -04  
  
25/168            1s 9ms/step - loss: 1.4099e -04  
  
31/168            1s 9ms/step - loss: 1.4012e -04  
  
37/168            1s 9ms/step - loss: 1.4049e -04  
  
43/168            1s 9ms/step - loss: 1.4056e -04  
  
49/168            1s 9ms/step - loss: 1.4080e -04  
  
55/168            1s 9ms/step - loss: 1.4157e -04  
  
61/168            0s 9ms/step - loss: 1.4254e -04  
  
67/168            0s 9ms/step - loss: 1.4300e -04  
  
73/168            0s 9ms/step - loss: 1.4334e -04  
  
79/168            0s 9ms/step - loss: 1.4361e -04  
  
85/168            0s 9ms/step - loss: 1.4373e -04

91/168	0s 9ms/step - loss: 1.4374e -04
97/168	0s 9ms/step - loss: 1.4372e -04
103/168	0s 9ms/step - loss: 1.4377e -04
109/168	0s 9ms/step - loss: 1.4400e -04
115/168	0s 9ms/step - loss: 1.4436e -04
121/168	0s 9ms/step - loss: 1.4486e -04
128/168	0s 9ms/step - loss: 1.4530e -04
134/168	0s 9ms/step - loss: 1.4556e -04
141/168	0s 9ms/step - loss: 1.4581e -04
148/168	0s 9ms/step - loss: 1.4604e -04
154/168	0s 9ms/step - loss: 1.4619e -04
160/168	0s 9ms/step - loss: 1.4638e -04
166/168	0s 9ms/step - loss: 1.4654e -04
168/168	2s 10ms/step - loss: 1.5080e -04 - val_loss: 6.8355e -05



```

#LSTM evaluations
y_pred_lstm = model.predict(X_test_lstm).flatten()

lstm_metrics = compute_regression_metrics(y_test_lstm, y_pred_lstm)
lstm_dir_acc = directional_accuracy(y_test_lstm, y_pred_lstm)

print("LSTM Metrics:", lstm_metrics)
print(f"Directional Accuracy: {lstm_dir_acc:.2%}")

1/56           7s 138ms/step

18/56          0s 3ms/step

36/56          0s 3ms/step

53/56          0s 3ms/step

56/56          0s 5ms/step

56/56          0s 6ms/step

LSTM Metrics: {'RMSE': np.float64(0.021962368050795756), 'MAE': 0.010625985694445502, 'MAPE': 46.98%}

summary_rows = []

summary_rows.append(
    {"Model": "Naive",
     **baseline_metrics,
     "DirAcc": baseline_dir_acc}
)

summary_rows.append(
    {"Model": "LinearRegression",
     **lr_test_metrics,
     "DirAcc": lr_dir_acc}
)

summary_rows.append(
    {"Model": "RandomForest",
     **rf_test_metrics,
     "DirAcc": rf_dir_acc}
)

```

```

    )

summary_rows.append(
    {"Model": "ARIMA(1,0,1)",
     **arima_metrics,
     "DirAcc": arima_dir_acc}
)

summary_rows.append(
    {"Model": "LSTM",
     **lstm_metrics,
     "DirAcc": lstm_dir_acc}
)

summary_df = pd.DataFrame(summary_rows)
summary_df.to_csv("../data/final_metrics.csv")
summary_df.head()

```

	Model	RMSE	MAE	MAPE	R2	DirAcc
0	Naive	1.694052e- 02	1.160951e- 02	5.502865e+07	0.509162	
1	LinearRegr	4.91b4564e- 07	1.411566e- 07	6.262626e+02	0.000000	1.000000
2	RandomForest	1.71947e- 03	8.215782e- 05	4.043183e+00	991546	0.999445
3	ARIMA(1,0,1)	0.88510e- 02	7.145005e- 03	3.433071e+07	0.000894	0.545656
4	LSTM	2.196237e- 02	1.062599e- 02	4.704331e+07	1.986357	0.469791

```

fig, axes = plt.subplots(1, 2, figsize=(12, 4))

axes[0].bar(summary_df["Model"], summary_df["R2"])
axes[0].set_title("R2 by Model")
axes[0].set_xticklabels(summary_df["Model"], rotation=45, ha="right")

axes[1].bar(summary_df["Model"], summary_df["DirAcc"])
axes[1].set_title("Directional Accuracy by Model")
axes[1].set_xticklabels(summary_df["Model"], rotation=45, ha="right")

plt.tight_layout()
plt.savefig("../images/R2_Accuracy.png")
plt.show()

```

```
/tmp/ipykernel_4825/602907317.py:5: UserWarning: set_ticklabels() should only be used with a
  axes[0].set_xticklabels(summary_df["Model"], rotation=45, ha="right")
/tmpp/ipykernel_4825/602907317.py:9: UserWarning: set_ticklabels() should only be used with a
  axes[1].set_xticklabels(summary_df["Model"], rotation=45, ha="right")
```

